

Intro to PCAP

Reid Gilman

Approved for Public Release: 13-0979. Distribution Unlimited

Creative Commons

This presentation is licensed under a

Creative Commons

Attribution-NonCommercial-ShareAlike 3.0

License: (

<http://creativecommons.org/licenses/by-sa/3.0/>)

<http://opensecuritytraining.info/>

This Is Not A Networking Class

It is about **problem solving**

Goals

1. Teach problem solving strategies using network analysis examples
2. Demystify the fundamentals
3. Know your tools

Syllabus

Day 1

1. Why PCAP?
2. Collection Techniques
3. PCAP Storage
4. Berkeley Packet Filter
5. Connectivity Problems
- 6. Lunch**
7. HTTP
8. Chopshop

Day 2

1. Unknown Protocols
2. DNS
- 3. Lunch**
4. Final Exercise

WHY PCAP?

(a.k.a. Who Cares?)

What Is PCAP?

- PCAP == **P**acket **C**apture
- Complete record of network activity
 - Layers 2 – 7
- Most common format is `libpcap`
 - Open-source
 - Available on *nix and Windows
 - C library, bindings in many languages
 - Others proprietary formats not covered

Pop Quiz!

Who Remembers The OSI Model?

7	• Application
6	• Presentation
5	• Session
4	• Transport
3	• Network
2	• Data Link
1	• Physical

Use Cases


- Identify rogue DHCP servers
- Search for evidence of malware activity
 - “malicious traffic”
- Follow insider threat’s footsteps
- Audit bandwidth usage
- Passive DNS resolution
- Monitor intrusions
- Test research hypothesis

Who Uses PCAP?

- **Researchers:** access to raw data
- **Administrators:** debug network problems
- **Analysts:** characterize malware activity
- **Incident Responders:** follow malware

You!

Collecting PCAP



Fair Warning


- Any plaintext protocol will be visible
- Careful what you log in to
- You'll be surprised what uses plaintext

Exercise


Use Wireshark and tcpdump to capture traffic while you ping `google.com`. What is in the ICMP Echo Request payload?

- Both tools installed in your VM
- `"ping google.com"`
- You will need to read the tcpdump man page
- `"man tcpdump"`

Aggregating Taps



Inline Devices



Naïve PCAP Storage

$$1\text{gbps} \times 3600 \times 24 = 86400 \text{ gigabits}$$

$$86400 \div 8 = 10800 \text{ gigabytes}$$

- Double that for full-duplex
- Storage can get expensive quickly

Packets Per Second


$$1\text{gbps} \div (64 \text{ bytes} \times 8 \text{ bits}) =$$

$$10^9 \text{ bits/s} \div 512 \text{ bits} = 1953125\text{pps}$$

Let H_i represent the overhead of storing one packet

$$N \text{ pps} \times H_{\text{PCAP}} = NH \text{ Bps}$$

libpcap Format



Exercise

How much overhead does libpcap incur storing packets for one hour on a saturated simplex 1gbps link with an average packet size of 1500 bytes?

libpcap overhead

Avg. Packet Size (bytes)	Packets Per Second	Overhead (MB/s)	Overhead (GB / day)
64	1,953,125	7.45	628.64
1500	83,333	0.32	26.82
7981	15,662	0.06	5.04
9000	13,888	0.05	4.47

Retention Policies

What to keep and for how long?

Data	Example Retention Period
Full PCAP	Weeks - Months
Flow Records	Indefinitely
DNS	Indefinitely
First N Bytes	Months - Years

BERKELEY PACKET FILTER

Surprisingly Powerful

Berkeley Packet Filter

- a.k.a. BPF
- “`man pcap-filter`” on Unix systems
- Conceptually similar to Wireshark filters
- Filter on layer 2+
- Richest in layers 2 – 4
- Very fast

Filtering Techniques

- BPF is limited, but fast
 - Compiles to an optimized form
 - Almost certainly faster than filters you write
- If you can use BPF, do it

Demo: Counting TCP Packets

You know a particular backdoor sends exactly one message per TCP packet.

How can you use tcpdump and command line tools to get a rough count of how many messages have been sent?

BPF Logic

- Combine BPF primitives with logical operators
 - NOT, AND, OR
- Easy to filter host and TCP/UDP port
- Advanced filters for TCP, UDP, ICMP, etc.
- Access to raw packet bytes

What Does This Do?

host 8.8.4.4 and udp port 53



Only traffic to
or from this IP



Only traffic to
or from this
UDP port

How About This?

```
dst host 74.125.228.36 and  
icmp[icmptype] = icmp-echo
```

How About This?

`ip dst 74.125.228.36 and`



Only traffic to this IP

`icmp[icmptype] = icmp-echo`



Filter on ICMP type

One More

```
ip[2:2] >= 86 and ip[8:1] <= 4  
and tcp[13:1] & 4 == 4
```

One More

`ip[2:2] >= 86 and ip[8:1] <= 4`



IP Length >= 86



IP TTL <= 4

`and tcp[13:1] & 4 == 4`



TCP RST

Exercise


You need to be notified immediately if anyone sets up a successful TCP handshake to 172.16.191.1 on TCP port 80 or if they send it more than 200 bytes on UDP port 53. Look at alert.pcap.

Write a script using tcpdump that will send you an email when either condition triggers.

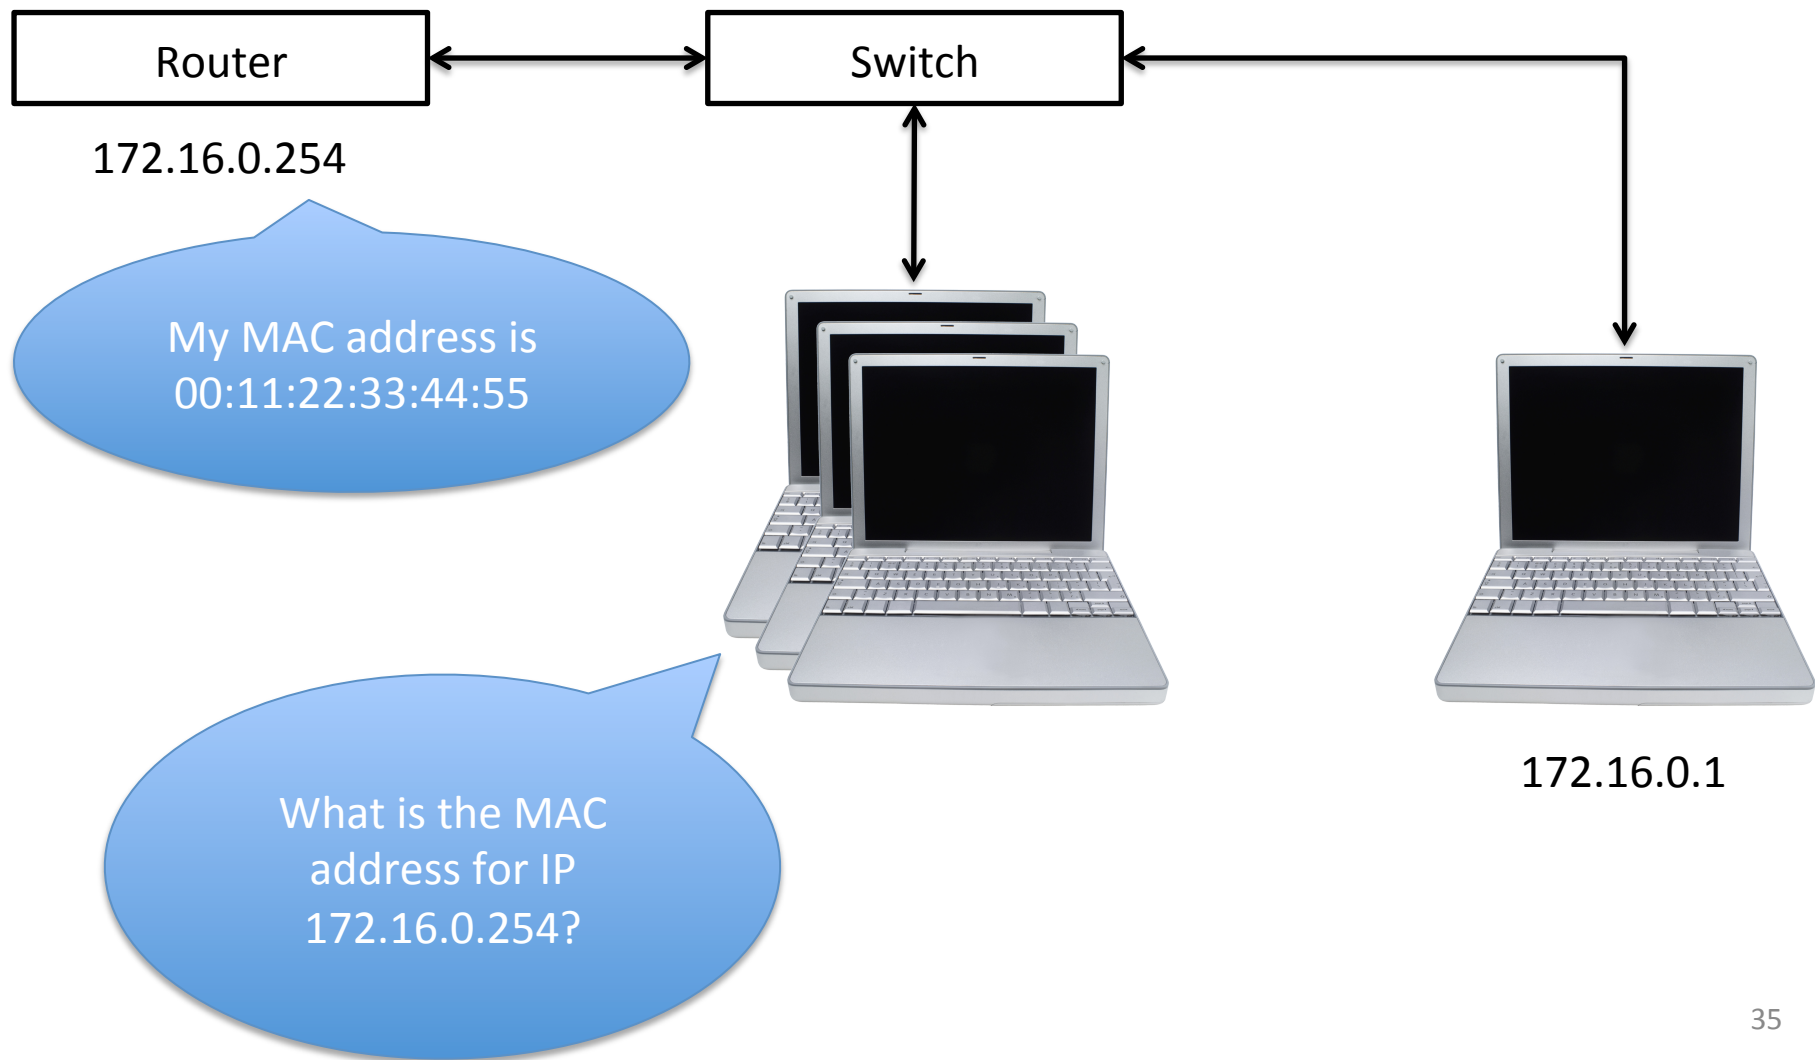
ADDRESS RESOLUTION PROTOCOL

Is It Plugged In?

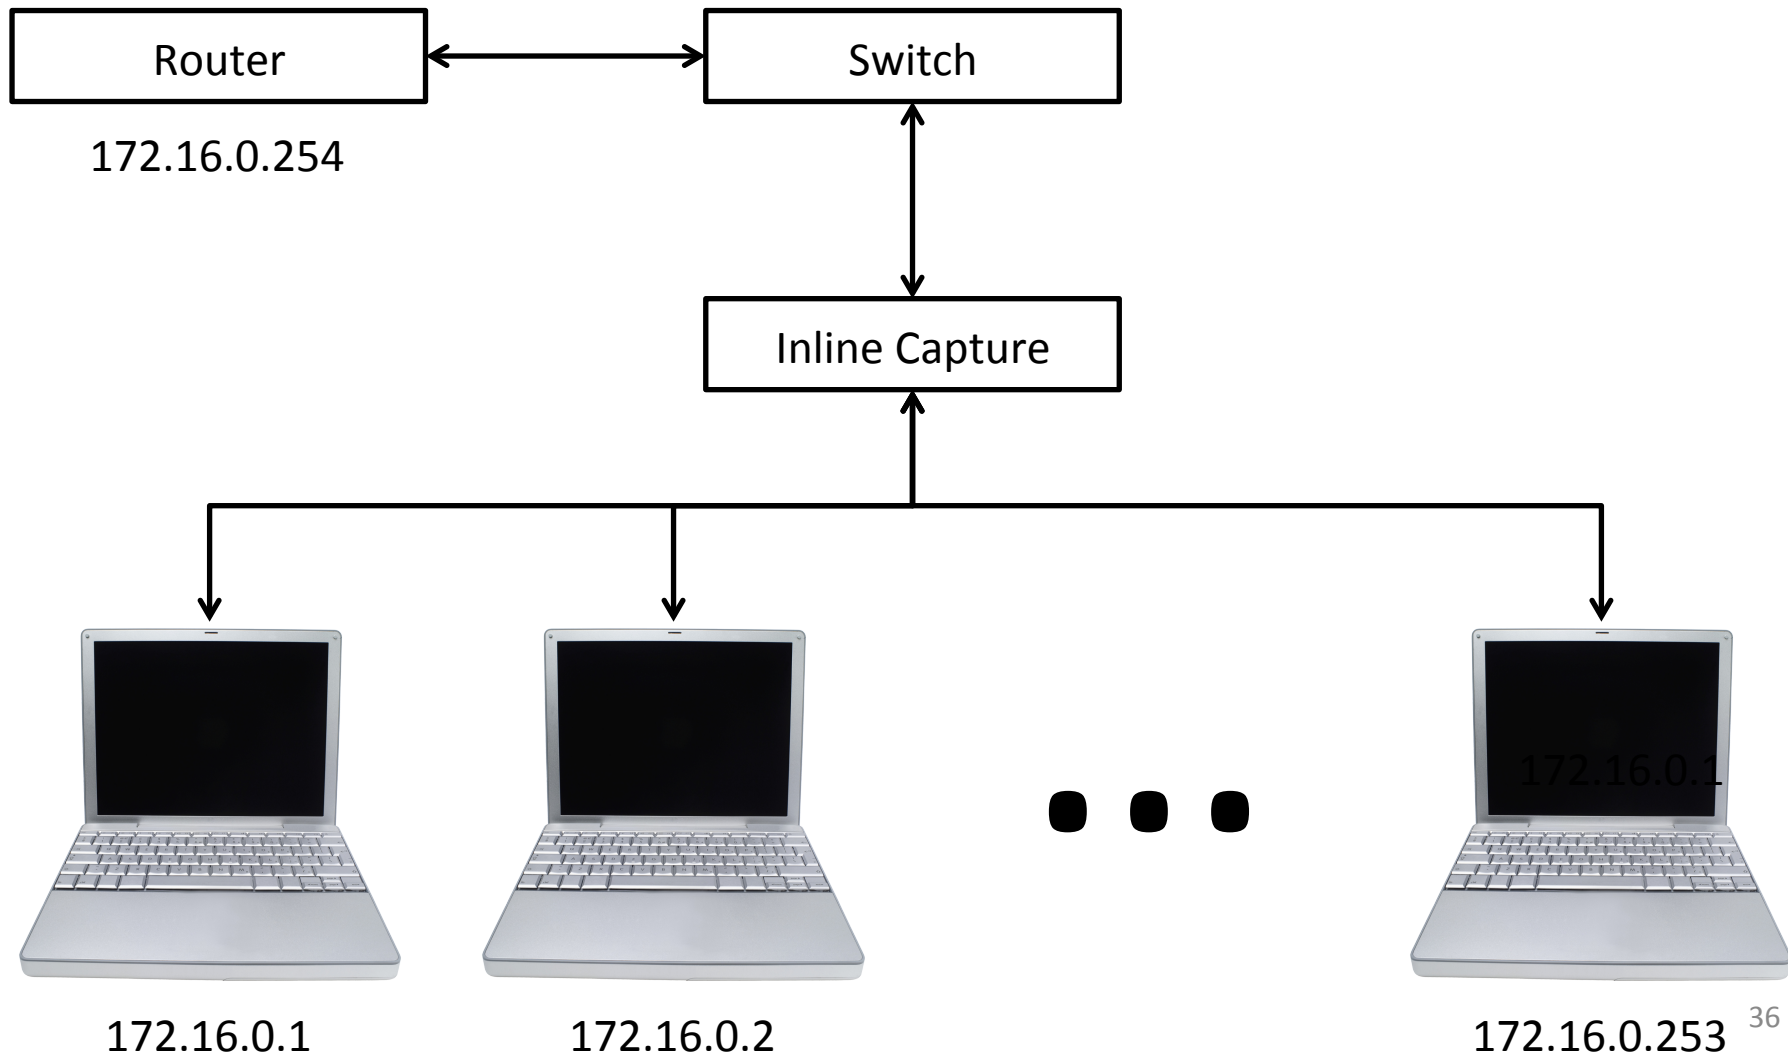
Network Connectivity



ARP



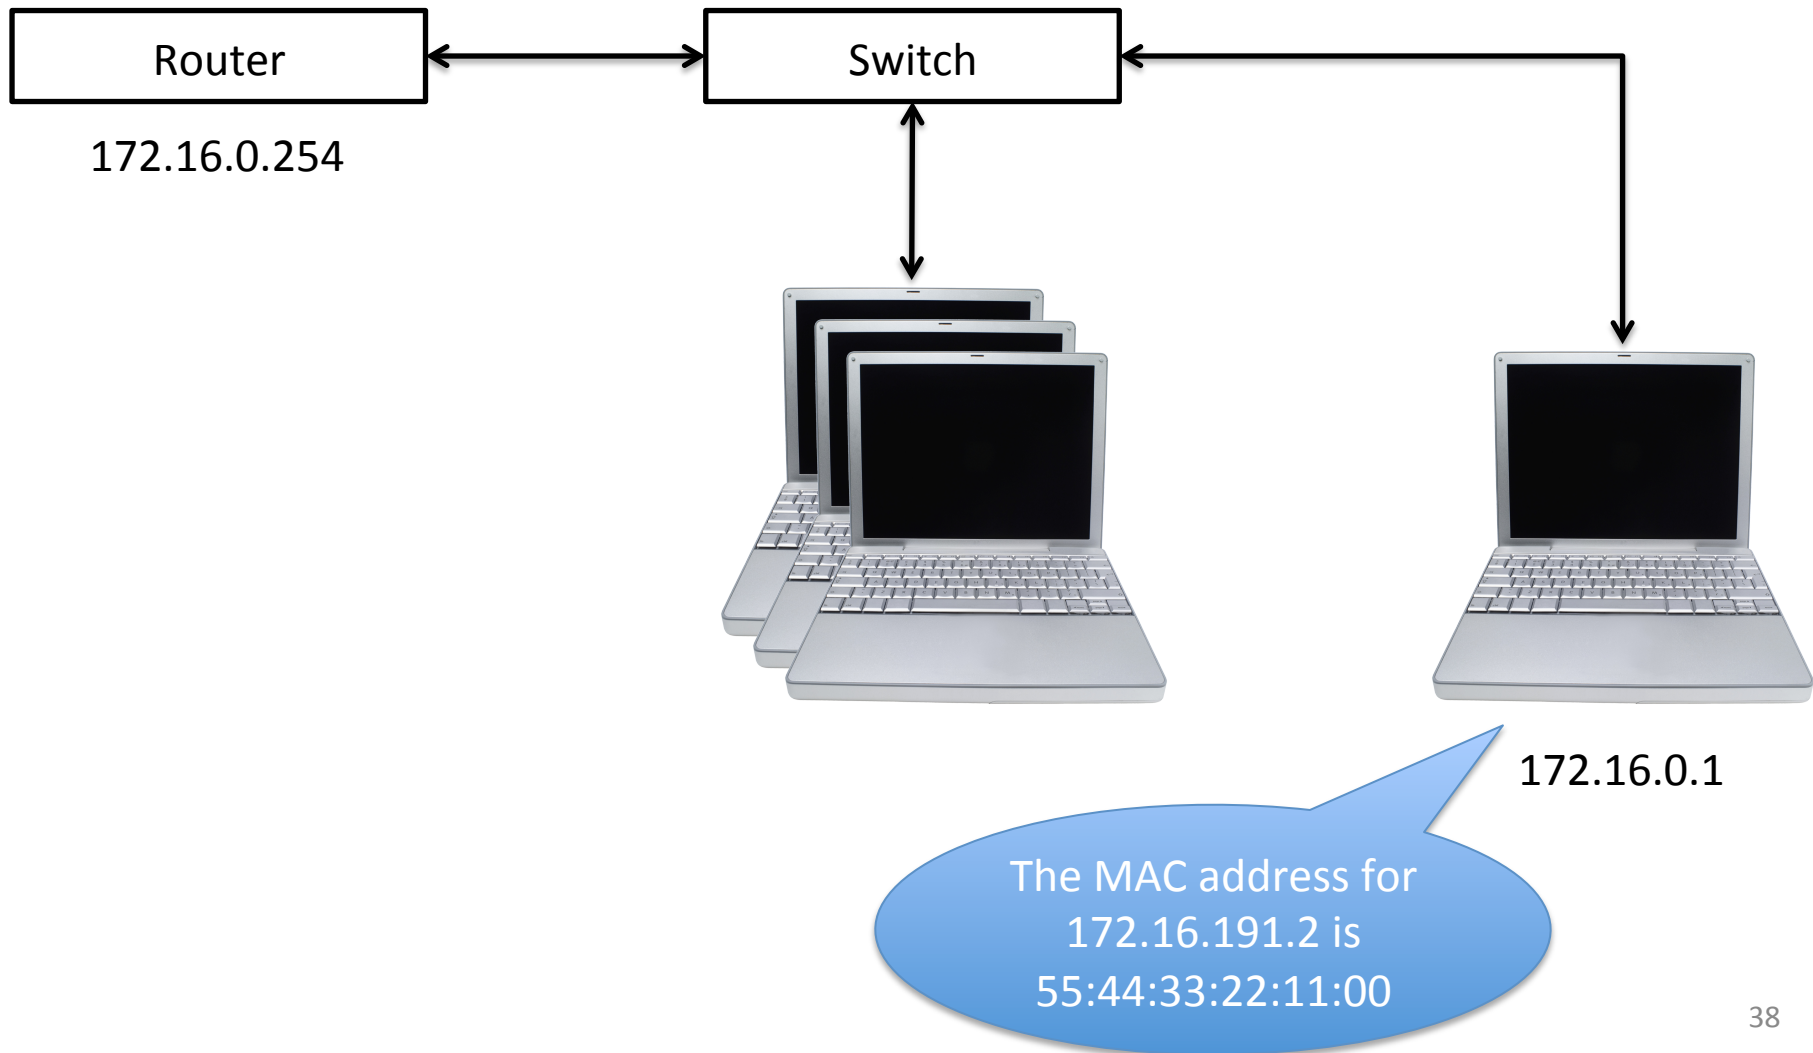
Exercise



Exercise

- Open `arp[0–9] { 3 } .pcap`
- `arpN.pcap` shows traffic from `172.16.0.N`
- Identify:
 - Default router IP address
 - Default router MAC address
 - IP and MAC address mappings

ARP Poisoning



ARP Poisoning

- Intercept all local traffic
- Low processor requirements
- Existing tools:
 - arpspoof + fragroute
 - sslstrip
 - Ettercap
 - Cain & Abel

HTTP

Hypertext Transfer Protocol

- Line-based protocol
- Intuitive fundamentals
- Many corner-cases

- Ubiquitous
- Many uses

Line-Based

- Headers are separated by line breaks
 - “\r\n”
 - Carriage-Return, Line-Feed
- Easy to read
- Works with existing line-based tools
 - grep, sed, awk, tr, etc.

HTTP

Header 1
Header 2
Header 3
Header n
Body

Headers

HTTP Verb

Protocol Version

GET / HTTP/1.0

Request Path

Headers

Header Value
↓
Host: www.google.com
↑
Header Name

Example

GET / HTTP/1.0

Host: www.google.com

User-Agent: wget

Connection: close

*BodyDataBodyDataBodyData
BodyDataBodyData*

Example

19db	5a85	e52c	af9c	4745	5420	6874	7470	..Z..,..GET.http
3a2f	2f78	6b63	642e	636f	6d2f	2048	5454	://xkcd.com/.HTT
502f	312e	310d	0a55	7365	722d	4167	656e	P/1.1..User-Agen
743a	2057	6765	742f	312e	3134	2028	6461	t:.Wget/1.14.(da
7277	696e	3131	2e34	2e32	290d	0a41	6363	rwin11.4.2)..Acc
6570	743a	202a	2f2a	0d0a	486f	7374	3a20	ept:.*/*..Host:.
786b	6364	2e63	6f6d	0d0a	436f	6e6e	6563	xkcd.com..Connec
7469	6f6e	3a20	436c	6f73	650d	0a50	726f	tion:.Close..Pro
7879	2d43	6f6e	6e65	6374	696f	6e3a	204b	xy-Connection:.K
6565	702d	416c	6976	650d	0a0d	0a		eep-Alive....

Example

19db	5a85	e52c	af9c	4745	5420	6874	7470	..Z..,..GET.http
3a2f	2f78	6b63	642e	636f	6d2f	2048	5454	://xkcd.com/.HTT
502f	312e	31 0d	0a 55	7365	722d	4167	656e	P/1.1..User-Agen
743a	2057	6765	742f	312e	3134	2028	6461	t:.Wget/1.14.(da
7277	696e	3131	2e34	2e32	29 0d	0a 41	6363	rwin11.4.2)..Acc
6570	743a	202a	2f2a	0d0a	486f	7374	3a20	ept:.*/*..Host:.
786b	6364	2e63	6f6d	0d0a	436f	6e6e	6563	xkcd.com..Connec
7469	6f6e	3a20	436c	6f73	65 0d	0a 50	726f	tion:.Close..Pro
7879	2d43	6f6e	6e65	6374	696f	6e3a	204b	xy-Connection:.K
6565	702d	416c	6976	65 0d	0a0d	0a		eep-Alive....



What's happening here?

Example

GET / HTTP/1.0\r\n ← CRLF splits headers

Host: www.google.com\r\n

User-Agent: wget\r\n

Connection: close\r\n

\r\n ← Blank line with CRLF ends headers

Body Data

Everybody Try This

```
$ echo -e "GET / HTTP/1.0\r\n  
> Host: www.google.com\r\n  
> \r\n" |  
> nc www.google.com 80
```

What Did That Do?

This Is Just Text

- How would you find a particular header?
 - It's value?
- Can you search for strings in the body?
- What is the response code?

Exercise: Find All The Titles

We need to extract all of the web page titles from a PCAP. Look in `http.pcap` for data. List every title exactly once.

Use `tcpdump` for this exercise.

Exercise: All Websites

Find all of the websites that the host 172.16.191.140 visited in `websites.pcap`. Do not list websites that other hosts visited. Don't forget about servers that may host multiple websites!

Limitations

- tcpdump and grep fall apart on large PCAPs
- tcpdump output not really parseable
- Could use libpcap or pynids
 - Lots of boilerplate code to get going
 - Not ideal for rapid prototyping

CHOPSHOP

Not Just For Cars

Chopshop

- <http://www.github.com/MITRECND/chopshop>
- MITRE-developed packet framework
 - Based on libnids
 - TCP reassembly
 - Handles boilerplate code
 - Python
 - Great for rapid prototyping

Chopshop

- Framework provides a standard API
- Framework does not analyze packets
- Modules provide all the brains
- Invoke with a list of PCAP files and modules

payloads

- Module to dump packet contents
- Useful for human-readable protocols
 - HTTP, SMTP, IMAP, etc.
- Few command line flags
- Can XOR data
- Can hexdump data
- Good first step in analysis

Invoking Chopshop

```
$ chopshop -f http.pcap "payloads "
```

– Run payloads module on http.pcap

```
$ find pcaps -type f |
```

```
> chopshop "payloads "
```

– Run payloads on all files in pcaps directory

Simple Obfuscation

- Many simple techniques are frustrating
 - Compression
 - Packing
 - Encoding
- Obfuscation is not encryption
 - No key required to “break” it
 - Still aggravating

XOR

- **Ex**clusive **O**r
- Basis for many ciphers
 - RC4, AES
- Fast in hardware
- Trivial in most programming languages
 - Typically a built-in operator
- Key management is easy

XOR Truth Table

Operand a	Operand b	Result
1	0	1
1	1	0
0	0	0
0	1	1

For $a, b \in \{0,1\}$, $a \oplus b$ is true iff $a \neq b$

In Other Words

One or the other,
but not both.

Examples

Operand <i>a</i>		Operand <i>b</i>		Result	
0111		1101		1010	
1100		1100		0000	
1100		1111		0011	
1111	0000	1100	1100	0011	1100
1011	0100	1010	1010	0001	1110
0x10		0x0A		0x1A	
0x10		0x32		0x22	

XORcise

```
$ chopshop -f xor.pcap "payloads "
```

- Add “-o 0xNN” to XOR contents
- Can you guess the XOR key?

File Carving From HTTP

- Demo: Wireshark
- Exercise: chopshop “http_extractor”

KNOWN UNKNOWN

Why Can't Everyone Use HTTP?

Unknown Protocols

- You will encounter something unfamiliar
 - Frequently without client or server to test
- Malware often uses custom protocols
- So do many proprietary programs

Can You Identify It?

- Well-known port?
- Unusual port/protocol pairing?
 - This will break Wireshark
- Constant or repeating values?
- Repeating structure or pattern?
 - Check beginning of packets
 - TLV is very common
 - <http://en.wikipedia.org/wiki/Type-length-value>

You Can't?

- Can you acquire a client or server?
 - May need to reverse it
- Ask around
- Consider obfuscation and encryption
- You may need to consider alternatives
 - Some things will always be a mystery

Endianness

Hexadecimal	Little Endian	Big Endian
10	1000	0010
432	3204	0432
5555	5555	5555
5432	3254	5432

- “network order” is big endian
- x86 is little endian

Warmup Exercise

1. Use tcpdump, chopshop, or Wireshark to identify the “unknown” flows in:
 - 1 . unknown-1 . pcap
 - 2 . unknown-2 . pcap
 - 3 . unknown-3 . pcap
2. Can you identify the traffic?
3. Can you decode it? How?

DNS

Domain Name System

Domain Name System

- Resolve names to IP addresses
 - e.g. www.google.com -> 74.125.228.3
- Most applications use DNS
- DNS servers configured in operating system
 - DHCP
 - /etc/resolv.conf
 - Windows NIC Configuration

Message Format

Header
Question
Answer
Authority
Additional

Header

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
ID															
QR	Opcode				AA	TC	RD	RA	Z			RCODE			
QDCOUNT															
ANCOUNT															
NSCOUNT															
ARCOUNT															

16 bits (2 bytes) wide

Reading Hexdumps

```
$tcpdump -i eth0 -XX -nn udp port 53
```

```
0x0000:  406c 8f4d 5c05 5057 a808 e000 0800 4500
0x0010:  0046 3783 0000 3c11 f4f4 8153 1472 8153
0x0020:  3b17 0035 dda1 0032 60e9 b150 8180 0001
0x0030:  0001 0000 0000 0478 6b63 6403 636f 6d00
0x0040:  0001 0001 c00c 0001 0001 0000 020b 0004
0x0050:  6b06 6a52
```

Exercise: DNS and tcpdump

Can you use a combination of tcpdump and grep to discover what IP address a name resolves to?

google.com

Look at dns-tcpdump.pcap

Solution

```
$tcpdump -r dns-tcpdump.pcap "udp  
port 53" | grep -A 1 google.com
```

dns_extractor

- Bundled chopshop module
- Examines UDP packets for DNS
- Prints or stores requests and responses

dns_extractor

```
$chopshop -f dns.pcap "dns_extractor -p"
```

- Print every DNS record

```
$chopshop -f dns.pcap -J out "dns_extractor -J"
```

- Write DNS records to file “out” in structured format

```
$find pcaps -type f |
```

```
> chopshop "dns_extractor -p"
```

- Run chopshop on every file in the “pcaps” directory

```
$find /example{01..03} -type f -iname > \*.pcap  
| sort | chopshop "dns_extractor -p"
```

- Examine PCAPs in three directories in sorted order

Exercise

List all DNS names and their resolutions found in `dns.pcap`. Only list each (name, resolution) pair once. You may use `chopshop` or `tcpdump`. If you can, do it with both.

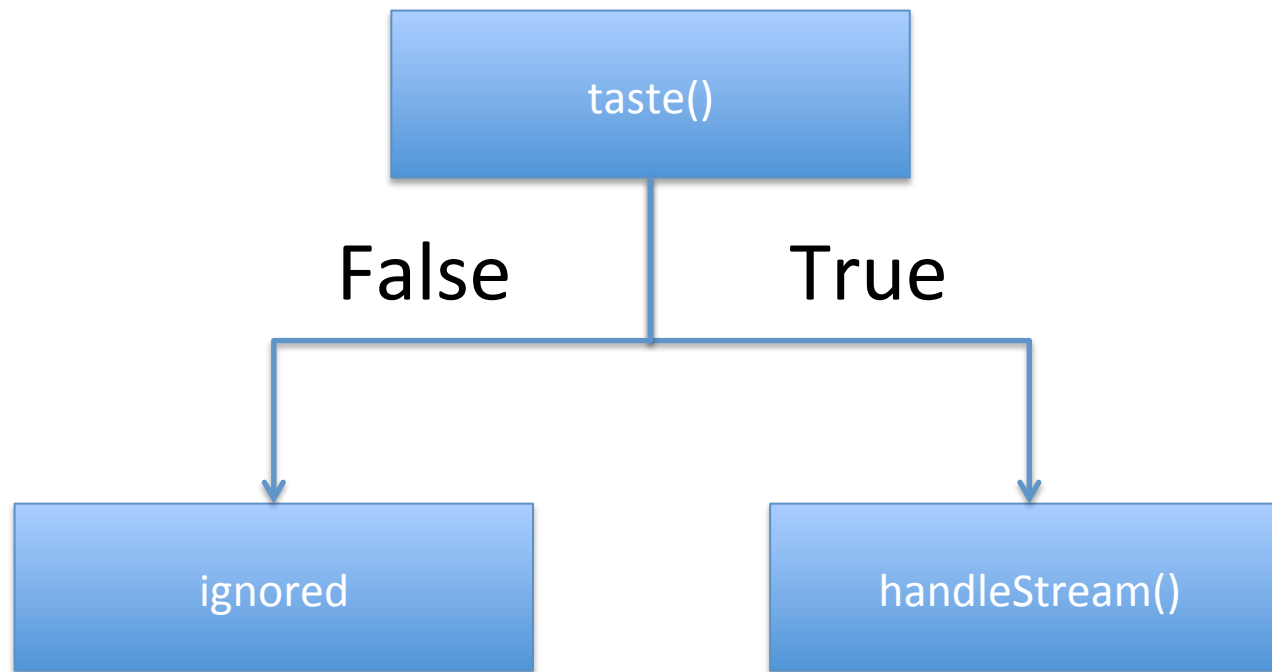
- Example “`file-reader.py`” might help
- Consider `chopshop`’s JSON-to-file option

Is Everything As It Seems?

- How many (name, resolution) pairs are there?
- How many DNS responses?
- How many `udp src port 53`?

CHOPSHOP MODULES

Module API Basics



Demo

- Open `my_first_module.py`

Mystery Port 53 Traffic

- How to understand this traffic?
- Is it encrypted or obfuscated?
 - With what algorithm?
 - Is there a key? Can you acquire it?

XOR in Python

```
s = "\x00\x01\x01\x01\x04\x05"
key = 0x01
out = ""
for char in s:
    out += chr(ord(char) ^ key)
```

XOR
Operator

Convert
character to int

XORcise

Decrypt and characterize the mystery traffic

Tips:

1. Remember the principles we just discussed
2. Iterate quickly
3. Worry about the process, not results

Outcomes

1. You captured PCAP
2. You worked on solving realistic, challenging PCAP analysis problems
3. You studied how and when to use different tools and how they might lie to you
4. You analyzed an unknown protocol

Hopefully...

1. You understand how PCAP can help you accomplish your goals
2. You can use a fundamental understanding of network protocols to go above and beyond existing tools
3. You had fun!